

Exploring Rural Shrink Smart Through Guided Discovery Dashboards

DENISE BRADFORD¹ AND SUSAN VANDERPLAS^{1,*}

¹*Department of Statistics, University of Nebraska - Lincoln, United States of America*

Abstract

Many small and rural places are shrinking. Interactive dashboards are the most common use cases for data visualization and context for exploratory data tools. In our paper, we will explore the specific scope of how dashboards are used in small and rural area to empower novice analysts to make data-driven decisions. Our framework will suggest a number of research directions to better support small and rural places from shrinking using an interactive dashboard design, implementation and use for the every day analyst.

Keywords *Interactive Dashboards; Exploratory Data Analysis (EDA); Parallel Coordinate Plots (PCP); Guided Discovery Learning (GDL).*

1 Abstract

Many small and rural places are shrinking. Interactive dashboards are the most common use cases for data visualization and context for exploratory data tools. In our paper, we will explore the specific scope of how dashboards are used in small and rural area to empower novice analysts to make data-driven decisions. Our framework will suggest a number of research directions to better support small and rural places from shrinking using an interactive dashboard design, implementation and use for the every day analyst.

2 Introduction

As the amount of data has increased in nearly every facet of life, the need to make sense of that data in an approachable, accessible form has become ever more important. As a result, many companies and organizations use interactive dashboards to present these data in a more useful and visually appealing form (Sarikaya et al., 2019). In many cases, dashboards support viewers' information processing, helping to make sense of complex data, navigate through a dataset, and supporting decision making based on the data. Dashboards are often used, as with the car display of the same name, to provide summary information about many separate attributes of a common entity. One glance at a car's dashboard will tell you the speed, RPM, engine temperature, amount of gas in the tank; more importantly, however, the goal is not for the user to remember all of these characteristics, but to assess whether any of these quantities is outside of the expected range. Similarly, interactive dashboards for data are often used to display many different attributes and performance metrics which are of importance for stakeholders. In this paper, we discuss the process of designing a dashboard to present publicly available government data to stakeholders

*Corresponding author Email: svanderplas@unl.edu.

in small Iowa towns to facilitate decision making and objective comparison with other similarly-situated towns.

Some communities continue to thrive as they lose population because they adapt, maintaining quality of life and community services for residents while investing in the future. This process, *smart shrinkage*, is important for rural areas who have experienced shrinking populations for decades. As small rural towns do not have access to data scientists or even the ability to easily leverage data collected locally to support decisions, our research team will provide communities with data about services in small town Iowa in order to assist with developing strategies to improve quality of life for their residents amid shrinking populations ([Rural Shrink Smart Team, 2022](#)). We hope to allow towns to explore their own data and compare to other similar towns, centering decision-making on data in the context of small-town Iowa life.

3 Data Description

The Smart and Connected Community (SCC) dashboard data are primarily assembled from [data.iowa.gov](#), with some additional datasets assembled from federal and private sources. Most of these data sets are collected at a town/city or county spatial resolution, requiring us to carefully join data to ensure that these differences are respected while collating relevant information at the city level. In addition to the more commonly available statistics derived from e.g. the census and American Community Survey, [data.iowa.gov](#) contains several unique data sets, including local liquor sales, school building locations, town budgets and expenditures, hospital beds, Medicaid reimbursements, and other details that may provide information about local quality of life.

Data available on Iowa’s data portal were augmented in some cases with higher-quality data sets in cases where the Iowa data were out of date or insufficiently accurate. Data collected from ELSI from <https://nces.ed.gov> were used to show the distance to any private or public school. The National Center for Education Statistics (NCES) is the primary federal entity for collecting and analyzing data related to education. [[@petersCommunityResiliencyDeclining2019](#)] – [Change the reference](#)

Data collected from the Index of Relative Rurality (IRR) ([USDA - ERS, 2020a](#)) were used in the SCC dashboard to help classify the towns. The Index of Relative Rurality (IRR) is a continuous, threshold-free, and unit-free measure of rurality. It is an alternative to the traditional discrete threshold-based classifications. The IRR ranges between 0 (low level of rurality, i.e., urban) and 1 (most rural). Four steps are involved in its design:

1. Identifying the dimensions of rurality: population size, density, remoteness, and built-up area.
2. Selecting measurable variables to adequately represent each dimension:
 - Size: logarithm of population size
 - Density: logarithm of population density.
 - Remoteness: network distance.
 - Built-up area: urban area (as defined by the US Census Bureau) as a percentage of total land area.
3. Re-scaling the variables onto bounded scales that range from 0 to 1.
4. Selecting a link function: unweighted average of the four re-scaled variable.

Data collected from Rural Urban Commuting Area Codes ([USDA - ERS, 2020b](#)) were used to help identify towns with commuting behaviors in our rural areas. The rural-urban commuting area (RUCA) codes classify U.S. census tracts using measures of population density, urbanization, and daily commuting. This data is on a zip code-level that will help identify those communities

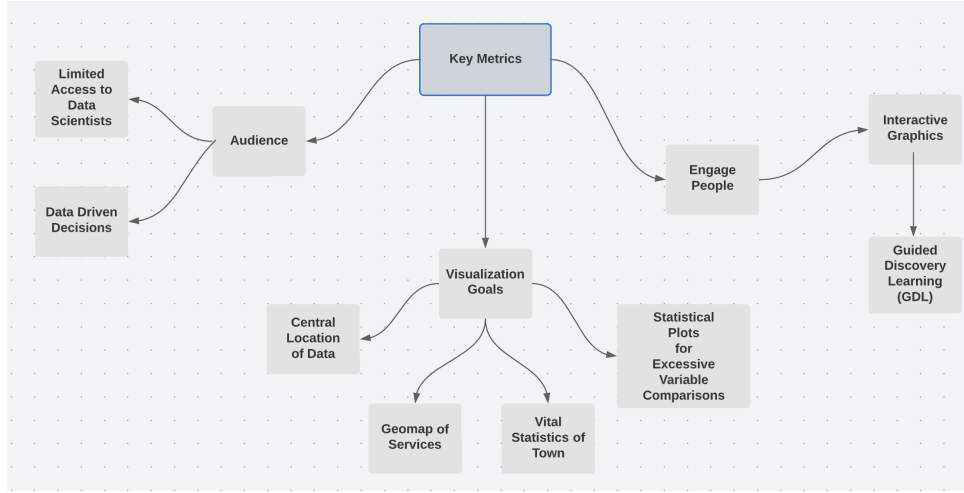


Figure 1: Diagram of considerations for our dashboard design process.

that commute to more urban areas. The most recent RUCA codes are based on data from the 2010 decennial census and the 2006-10 American Community Survey. The classification contains two levels. Whole numbers (1-10) delineate metropolitan, micropolitan, small town, and rural commuting areas based on the size and direction of the primary (largest) commuting flows.

One of the interesting features of this assembled data set is that missing data can be missing for multiple reasons: not all state data is complete, but data about certain services may also be missing because towns do not offer that service. Thus, in addition to the usual challenges of working with real-world data that is “messy” in a variety of ways, we also have to contend with missing data that is missing due to the size of the community or the lack of services. This makes both visualization and statistical analysis more complicated (and more interesting).

4 Dashboard Design Considerations

One problem we identified early in the process of assessing smart-shrinkage strategies in small towns is that these towns do not have the resources to make data-driven decisions. Typically, small towns in Iowa are managed by at most a few part-time employees or volunteers. In some cases, essential management functions of the town are paid, but the municipalities we are interested in do not have sufficient funding to hire professionals to gather and analyze data.

As part of a wider project investigating the strategies towns use to maintain quality of life amid shrinking population, our research team provides communities with data about their own town, but also comparable towns across the state which may have a different approach to city services. In combination with other engagement strategies that are more qualitative, we hope to use this interactive dashboard approach to assist small Iowa cities with generalizing and developing strategies to improve or maintain quality of life amid shrinking populations.

One factor at the forefront of our visualization design is the importance of reducing the cognitive demands on viewers: we have assembled an incredible amount of data, and it is easy for even statisticians who deal with much larger datasets to get lost in the details of this data. At the same time, we want to invite viewers to engage with the data - to imagine, to draw comparisons, to generalize across towns, and to integrate outside information into the conclusions

drawn based on the data we present. This invitation to engage with the data is similar to the approach advocated in Guided Discovery Learning, a framework leverages hints, feedback, and other helpful information to guide users in interactive exploration (DeDonno, 2016).

We expect that users will be interested in “sets” of variables from the wider dataset, which we assembled based on quality of life factors in the Iowa Small Town Poll (Peters, 2019). For instance, users might be interested in medical and social services available to residents, such as a local primary care clinic, nursing homes which are within driving distance, and the distance to the nearest emergency room; these factors might be explored separately from variables describing the services provided directly by city government, such as parks and recreation expenditures, snow removal services, and the distance to the closest fire station.

As a consequence of this massively multivariate structure, we very quickly focused on the use of parallel coordinate plots; other alternatives, such as tours (Wickham et al., 2011), require much more sustained attention to interactive plots as well as a deeper understanding of projections in multidimensional space which we cannot assume our users will have. Introduced in the 1880s (d’Ocagne, 1885), parallel coordinate or parallel set plots feature a series of vertical axes representing different variables arranged horizontally, with lines connecting each observation. When representing categorical data, parallel set plots may show “blocks” of data instead of individual lines, and are useful for representing conditional relationships between adjacent variables (Bendix et al., 2005); modifications of this design, such as common-angle plots (Hofmann and Vendettuoli, 2013), address the issues which arise due to line-width illusions (Day and Stecher, 1991, VanderPlas and Hofmann (2015)). Parallel coordinate plots have been generalized to allow for continuous data and additional summaries beyond individual data points, such as densities (Heinrich and Weiskopf, 2009). In this paper, we use the `ggpcp` package, which leverages the grammar-of-graphics framework introduced in Wickham (2016), allowing us to use not only parallel coordinate plots, but also to overlay other statistical summaries, such as boxplots or violin plots, to provide additional context about the marginal distributions of each variable in addition to allowing for exploration of the multivariate space.

We also anticipate that users will be interested in comparing their town to other, similar towns. We will discuss the different ways that this comparison strategy was implemented in each dashboard in the next section, which describes the evolution of the dashboard over time and accounting for feedback from users and other researchers on the wider project.

One final component of this project is that our dashboard is part of a wider effort to work with towns to understand the different strategies used to maintain resident quality of life amid shrinking populations. Thus, while the town leaders are our primary audience, we also are creating this applet for use in parallel with a team of other researchers: sociologists, economists, city planning specialists, and artists. These researchers opinions and feedback about the dashboard are also useful and important, as they regularly work with town leaders in different capacities and have an understanding of what factors are most important to them and what types of questions these leaders may have when faced with data and unfamiliar statistical visualizations.

Throughout the design process, we will assess our visualizations to determine which strategies for user interface and interactive graphics design are most useful to empower town leaders to make discoveries in publicly available data assembled with a focus on items that impact rural quality of life.

4.1 Guiding Design Principles

Research on dashboard creation and interactive visualization tends to be very task-specific and hard to apply to more generalized settings. That is, it is relatively easy to create a dashboard that works for a particular task, but it is hard to generalize from that process what will work for the next dashboard. With this in mind, we set out to clearly document our intentions at each stage of the design and evaluation process, with the goal of gathering some useful information about general dashboard design from the process of creating this specific dashboard.

Thus, our initial set of dashboard design principles is as follows:

- The town leaders are the focus audience; thus, the town itself should be the central focus of the app.
- We should facilitate comparisons with other towns in order to allow the user to explore other potential solutions to offering services that enhance resident quality of life.
- We will present the user with peer comparisons in order to widen the scope of exploration beyond the initial set of obvious peers in the local region.
- We will implement feedback mechanisms that allow us to provide more detailed data and respond to feature requests to improve the dashboard design over time.

As with many dashboards, this project is under continuous development; while it makes for an unsatisfactory conclusion, we do not have a “final” dashboard design because the application will continue to evolve. However, we have some useful insights into the process of creating an application designed to invite users to explore a large and complex dataset that we believe to be a useful contribution to work in this area.

5 Dashboard Design Process

5.1 Dashboard Components

In this section, we discuss the philosophy behind the basic “building blocks” of the dashboard. This philosophy is present in all of the iterations of the dashboard that we present in this discussion, and we will evaluate the overall philosophy’s effectiveness in the conclusion.

The large set of publicly available data (primarily from data.iowa.gov) we have assembled is useful, but we must be careful with how we present this data because it would be easy to overwhelm the user with small details that mask the bigger picture. We select a small subset of towns (out of the 999 towns in Iowa) and a small subset of variables of interest to start with, and then allow the user to increase the complexity of the display in accordance with their interest. This avoids some of the pitfalls of dashboard design that can easily lead to user overload (Few, 2006).

Our primary objective is to provide users with a town-centric approach: their town is at the center of our application, and comparisons to other, similar towns are secondary. As a result, the next component of the dashboard is intended to provide a brief overview of the information we have about a specific town of interest. This design is based on research into visualization sensemaking (Lee et al., 2016), in that we allow users to explore outward from the familiar to the unknown. The map visuals were built using Open Source Routing Machine (OSRM) route functions [luxen-vetter-2011] in R [r] to amplify the accuracy of the distances from necessary services in town-centric point. OSRM allows for finding the "As the Crow Flies" distance and time on the road for our vital services map, since OSRM technology is similar to Google maps.

When faced with the next component, a parallel coordinate plot (PCP), a novice user will be

able to determine two basic components: Visual Object (textual objects and non-textual objects) and Frame (frame of content and frame of visual encoding).

Taken together, the app is a single page; the initial “solid ground” which the user explores from consists of maps showing the route from the center of town to necessary services, including the fire department, schools, post offices, and hospitals. In version 2, as shown in Figure 3, the map portion is condensed, and more space is given to value boxes that show vital statistics about the town’s QoL and financial metrics. This relatively straightforward display is followed by a parallel coordinate plot that allows the user to see similar towns along dimensions such as economic indicators or population size.

5.2 Initial Draft

The initial design sketch and implementation are shown in Figure 2.

Users’ towns **Town Users’ ??** are at the center of our application, and comparisons to other, similar towns are secondary. As it can be extremely difficult to predict which towns are optimal for comparison purposes (similar may involve population, region, economic indicators, sports rivalries, and any number of other variables), we allow users to modify a set of suggested comparison towns to indicate other towns of interest.

We implemented some suggested town comparisons using unsupervised clustering methods to help our towns make decisions that are informed in comparison to similar towns, for budget size, population size and location. We initially focused on determining the next five to ten similar towns, based on distances to services. This feature became an important diagnostic for our data quality, as it became clear that towns which were grouped with big cities but which did not have a large population were so grouped because of missing data. Unfortunately, this clustering feature was not as useful to the application users, as they came to the dashboard with a pre-existing set of towns to compare to; our suggested comparisons were in the way.

The initial dashboard design featured several responsive maps showing the distance to the nearest hospital, fire department, post office, and school. These maps were ineffective for several reasons:

- Town residents already know this information (though it was useful for us as the dashboard designers, because we aren’t nearly as familiar with the 900+ small towns in Iowa)
- We computed distance from services relative to the center of town - coordinates provided in the data from data.iowa.gov. Generally speaking, the post office is at the center of town and the fire department is usually very close to the center of town; these two maps were useless. The school and hospital maps were less useless, but still did not provide particularly useful information to people already familiar with the town.
- It became clear that it might be more useful to show the comparison towns on a map (relative to the town of interest) so that users could compare geographical ratings for unfamiliar data to familiar data.

In addition, we received feedback on the parallel coordinate plot at the bottom of the app which was surprising: the viewers (in this case, other researchers on the team) were not as intimidated by the parallel coordinate plot as we had expected. They did need some explanation of how to read the plot, and these hints need to be included in the dashboard, but they grasped the fundamental idea of the plot very quickly.

Our conclusion, based on this initial dashboard draft, was that we needed to restructure the application. Our attempt to show familiar information first to “build up” to the more unfamiliar structure of a parallel coordinate plot was not effective; there was too much clutter and not

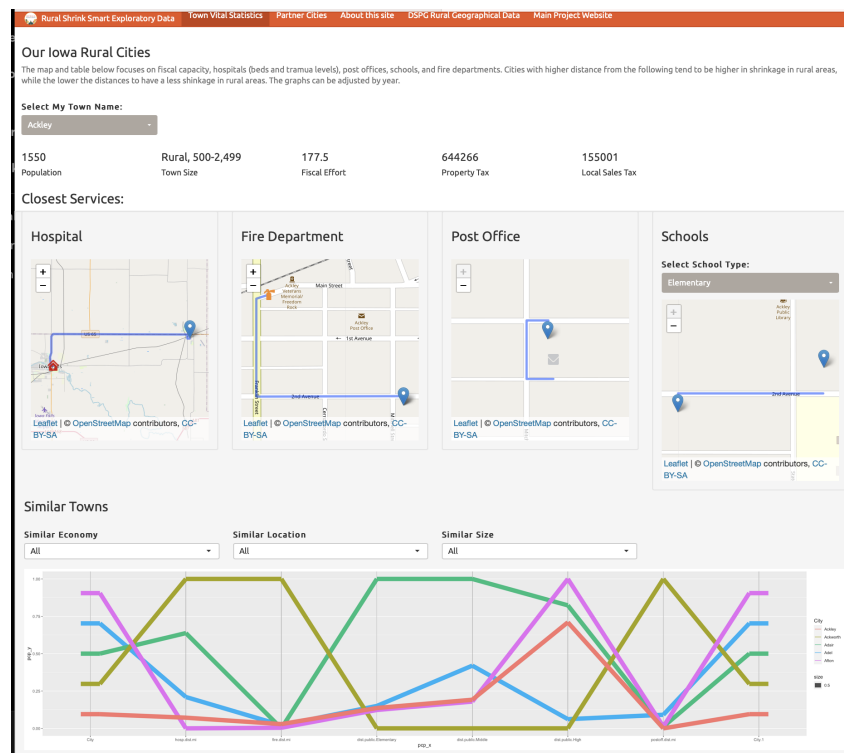


Figure 2: Initial dashboard design sketch (top) and implementation (bottom).

enough new information to draw users in.

5.3 Redesign

In the initial design, we included a map for each vital service, this initially created a lag for the users' experience. As a result, we cached map directions from OSRM for each service in our database, which drastically reduced the response time for the user. Our initial design did not naturally focus the user's eye on the most important parts of the dashboard; the redesign allowed for a cleaner flow from the top to the bottom.

In addition to the timing due to the map loading slowly, we added the vital statistics at the county level to allow for a more robust understanding of the town and its surroundings. The rurality index provided a better classification and the USDA sources allowed for the town to understand the impact of the closest major city due to commuting for work and shopping at larger stores not available within the town.

We also modified the parallel coordinate plots in several ways:

- Our x-axis had a large number of variables that we as researchers believed to be the most strongly associated with quality of life. However, there were still too many variables for users to successfully parse. We reduced the number of variables, focusing on variables that had the highest data quality, and we grouped these variables by quality of life factor (Peters, 2019).
- Originally, parallel coordinate bands were scaled based on the selected comparison towns. This had the effect of truncating the range of variables and over-emphasizing differences between selected towns relative to the overall range of each variable over all towns in our data set. We chose to show all towns in the data set in a very light α grey color to provide some information about the overall range of each variable. Unfortunately, even with the low- α value, this increased the visual complexity of the plot and confused users. Future iterations will likely make use of another aesthetic, such as boxplots or violin plots, to show the range of values for all towns, and then use lines only for towns that are selected by the user. This should strike a balance between visual complexity and representing the data accurately.
- We noticed that users did not make use of our suggested comparison towns, and so we removed that option in favor of allowing users to enter their own comparison towns directly. Users already had pre-determined towns they wanted to compare to, and our suggestions were just in the way.

While not all of these modifications were well received in our second round of user testing, the changes did incrementally move the dashboard display towards our goal of allowing users to explore the data and engage with it. We continued to be surprised with how well users reacted to the parallel coordinate plots, which we initially thought might be too abstract for users unfamiliar with multivariate data displays, but the ability to compare towns across multiple dimensions and examine the similarities and differences between their approaches to different services seemed to be intuitive for users once they understood that each vertical axis was a different variable.

6 Discussion

How did the philosophy work out?

The design philosophy worked well in the ways to promote guided discovery learning for the user. The blocks that were created helped create a way for the user to scroll to the next section. This framework in the design allowed for the highlighted pieces of the town in generally. This design allowed the user to see the PCP after understanding how the dashboard allowed for their town



Figure 3: A second iteration of the sketched design (top) and the implementation (bottom).

to centered. The PCP in the lower part of the dashboard allowed for the user to see the plot and adjust to the fact that they could add more towns to the plot.

What worked and didn't (overall)? The idea to use the guided discovery methodology created a lack of understanding in the plot for a novice user. The need to explain a visual means that the user is more likely to not use the plot effectively. Our hope was that the PCP would be a bit more "self-explanatory" than it was initially. The number of helpful hints clutters the UI and may be a distraction for a particular user.

People didn't seem to like having their community facts to explore from - they already knew that information - so what should we replace it with? Collaborator suggestions included reporting capability and data stories (narratives from the communities), but neither one really inspires exploratory data analysis. What we need to do is provide enough interactivity to allow people to select interesting data sets and use that to compare

Many small towns don't like to see that their information is being displayed publicly or used to be compared to another city. This could a huge barrier for our resarch and work. To make sure that this does not become a problem, we could make sure that when allowing the user to compare towns in the PCP, we make sure to de-identitfy the towns in the UI. For example, the town user selects similiar population size, the each town that is populated will have a collection of ten similiar towns using a label, such as "Town 1", "Town 2", etc. Along with the de-identifying other towns within the UI, using a login system for each town will allow for the town to feel as if their data is being protected. This login system can then be used as a filter to setup the dashboard there forward. This will hopefully remove the lack of ability to be interactive. This will seem as if the dashbaord is developed for their needs particularly. Once, the privacy issue is removed to understand more of the data, we should allow for the user to add/remove variables in to the plot as they see fit. Suggesting that the QoL variables are the ones that are added in the start, but allowing the user to remove subjects from the QoL sections will allow for the user to not only feel as if the UI is interactive, it would also create more exploration for the user.

We also should provide much, much more user support for how to interact with and interpret the app and results. Support for the dashboard should have a help email for towns to feel supported with interpreting the app and results. This helpful email will be connected to our research group. In which the help will not only have statisticians but also city planners that can help answer questions. As town users are looking to make changes for their towns, it could be a QoL variable that will change the status of the town, which can be addressed with the statisticians. We will be able to help them with making sure that the right variables are being used in the PCPs.

7 Future Work

Incorporating other dashboards which had slightly different focuses (99 vs 990 towns, general data display, etc.). In our research, our future work will incorporate other dashboards with general information that are slightly different focuses (99 vs 990 towns, general data display, etc.). In collaboration with DSPG at Iowa State University and the statistical modeling portions of our NSF group, we would like to transition our app to a multi-tab dashboard which allows for more specific task-oriented exploration and evaluation, rather than broad EDA and discovery.

Discussion of the difficulty of working with researchers from other disciplines, not because of convention but because they don't know what is possible with EDA and are unfamiliar with all but goal-directed data analysis.

EDA can be a difficult for a variety of groups of people, novice users and experienced researchers. While working with a variety of disciplines because EDA can be hard to understand the use cases of exploration of data. Many experience researchers, see EDA as a way to look at data that were collected to only describe the behavior of variables that are added into models or research questions. However, EDA can be used in a data discovery, such as the work displayed in this article. Data Discovery can help with the quality of the data, in order to find new resources. EDA can help with potential causality of town related issues. For example, if data are showing that people are no longer using public transportation, this could be due to the population shrinking but exploring the data could suggest that more people are using ride share companies. This could as suggest that public transportation may not offer the needs for the people in the town. EDA will help identify those issues or even allow for the town leaders to look for answers.

8 Conclusions

Specific takeaways - visualization research (broad) vs. dashboard design (very problem specific) is missing a critical set of research describing the specific and trying to make general suggestions as to how to proceed with the design process. [In conclusion](#),

Helpful hints - showing a collection of very different apps to train collaborators on what is possible works well to empower collaborators to help brainstorm. Test early, test often. . .

8.1 User Testing

In this paper, we present the results from testing this application with users, and examine the implications for the feedback we received on our application design and on the more general design of dashboards for statistical novices. We examine which aspects of dashboard design “common knowledge” hold up in practice, and assess why these principles may not hold in our use-case.

Complex visualizations are difficult to assess in vivo: it is fairly easy to design an experiment to test a single graphic, but much harder to test the interactive dashboards which consist of systems of displays and user interactions. As a result, it is important to continuously assess which “common knowledge” design principles stemming from more controlled experiments hold up in real-world situations.

Using the framework from “Creating design guidelines for building city dashboards from a user’s perspectives” ([Lee et al., 2016](#)), we will help summarize a users concerns utilizing six main critiques, which are framed into questions. We will adapt these questions to better help the town leaders communicate clear feedback for our team to use. We will collect how the users are experiencing the app with these set questions:

- Is the information the dashboard presents clear? (*Scope*)
- Is the dashboard easy to navigate? (*Usability*)
- Is the dashboard used in a socially responsible manner? (*Ethics*)
- Do you believe that you can trust the Rural Shrink Smart dashboard? (*Validity*)
- Do you find value in the Rural Shrink Smart dashboard? (*Utility*)

8.1.1 Beta Testing

Our Beta testing of a design dashboard before deploying the product to our town leaders, we allowed users from the following groups:

- Graphics Group Members (Users know Statistical Graphics)

- Rural Shrink Smart Project Members (Users know the project with little background in Statistical Graphics)
- Others (Member of doctoral community that has background in educational adaptivity)

While this may not be a large sample size, our range of testers helped with changing approaches in data collection and dashboard changes.

9 Discussion and Future Research

9.1 Data Collection Revisited

As mentioned above, we wanted to use Unsupervised Clustering Methods to determine similar towns, however we were able to use this methodology to determine data quality in our open data sources. This work empowered our time to look for other open data sources to validate our vital statistics and maps. Since the clustering methods were used differently, we decided to use the USDA's version of rural area definitions. This allowed for a more transparent look into towns that may be similar based on their rurality index. Similar to the Census, the rurality metrics are county based and updated every ten years. This could later cause a problem for towns that may be shrinking faster than the collection of data. It could be worth looking into developing a proxy from our data collected that would be reliable for the years in between the USDA's collected years.

In a related idea of making sure that the towns are identified correctly, we should look into the aging of our towns. This may be important to identify the reasons for shrinking. In the event, that the towns are aging, will we be able to identify how fast the younger generation is moving out of the towns. This could be the most important piece of truly identifying the use of data in towns with an aging population.

9.2 Visualization Revisited

In our statistical graphs, we used Parallel Coordinate Plots (PCPs) to allow the user to look at a variety of variables, while the Guided Discovery Learning was used to make sure that the users were informed of the information that was presented, it did not allow for the town leaders any more ease in understanding. As a result, we decided to look for easier ways to display data, such as boxplots and barplots. This direction will be limited but may be useful for the sake of ease on the town leaders.

In the event that we want to continue to have the town leaders explore using a different statistical graphic, we would like to introduce the flower plot. This method may not represent that data in the way that could be misleading for understanding, we believe that the flower plot will allow for town leaders to engage in being curious as they explore more dimensions of data related to their Quality of Life (QoL) survey.

9.3 Feedback

We will continue to engage with town leaders to ensure that we are developing a product that will help them make better data-driven decisions that are useful to increasing their quality of life, which may lend itself to help less shrinkage in these rural areas.

References

- Bendix F, Kosara R, Hauser H (2005). Parallel Sets: Visual Analysis of Categorical Data. In: *2005 IEEE Symposium on Information Visualization (INFOVIS'05)*, 133–140. IEEE.
- Day RH, Stecher EJ (1991). Sine of an illusion. *Perception*, 20: 49–55.
- DeDonno MA (2016). The influence of iq on pure discovery and guided discovery learning of a complex real-world task. *Learning and Individual Differences*, 49: 11–16.
- d'Ocagne M (1885). Coordonnées parallèles et axiales : Méthode de transformation géométrique et procédé nouveau de calcul graphique déduits de la considération des coordonnées parallèles. *Gauthier-Villars*, 112.
- Few S (2006). *Information dashboard design: The effective visual communication of data*. O'Reilly Media, Inc., Newton, MA.
- Heinrich J, Weiskopf D (2009). Continuous Parallel Coordinates. *IEEE Transactions on Visualization and Computer Graphics*, 15(6): 1531–1538.
- Hofmann H, Vendettuoli M (2013). Common Angle Plots as Perception-True Visualizations of Categorical Associations. *IEEE Transactions on Visualization and Computer Graphics*, 19(12): 2297–2305.
- Lee S, Kim SH, Hung YH, Lam H, ah Kang Y, Yi JS (2016). How do people make sense of unfamiliar visualizations?: A grounded model of novice's information visualization sensemaking. *IEEE*, 22: 499–508.
- Peters DJ (2019). Community Resiliency in Declining Small Towns: Impact of Population Loss on Quality of Life over 20 Years. *Rural Sociology*, 84(4): 635–668.
- Rural Shrink Smart Team (2022). Rural shrink smart. <https://ruralshrinksmart.org/>. Accessed: 2021-12-01.
- Sarikaya A, Correll M, Bartram L, Tory M, Fisher D (2019). What do we talk about when we talk about dashboards? *IEEE Transactions on Visualization and Computer Graphics*, 25(1): 682–692.
- USDA - ERS (2020a). Rural classifications. <https://www.ers.usda.gov/topics/rural-economy-population/rural-classifications/>. Accessed: 2022-04-01.
- USDA - ERS (2020b). Rural-urban commuting area codes. <https://www.ers.usda.gov/data-products/rural-urban-commuting-area-codes/>. Accessed: 2022-04-01.
- VanderPlas S, Hofmann H (2015). Signs of the sine illusion—why we need to care. *Journal of Computational and Graphical Statistics*, 24(4): 1170–1190.
- Wickham H (2016). *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York, 2 edition.
- Wickham H, Cook D, Hofmann H, Buja A (2011). tourr: An R Package for Exploring Multivariate Data with Projections. *Journal of Statistical Software, Articles*, 40(2): 1–18.